

# hearttrainaf

May 1, 2025

```
[3]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[4]: df=pd.read_csv('heart.csv')
```

```
[5]: df=df.drop_duplicates()
```

```
[6]: df.isna().sum()
```

```
[6]: id          0
age           0
sex           0
dataset       0
cp            0
trestbps     59
chol         30
fbs          90
restecg       2
thalch       55
exang        55
oldpeak      62
slope       309
ca          611
thal       486
num          0
dtype: int64
```

```
[7]: mean_col=['trestbps','chol','thalch','oldpeak','ca']
median_col=['fbs','restecg','exang','slope','thal']
```

```
[8]: for col in mean_col:
    if col in df.columns:
        df[col]=df[col].fillna(df[col].mean())

for col in median_col:
```

```
if col in df.columns:
    df[col]=df[col].fillna(df[col].mode()[0])
```

C:\Users\AMOL\AppData\Local\Temp\ipykernel\_9720\1465429545.py:7: FutureWarning: Downcasting object dtype arrays on .fillna, .ffill, .bfill is deprecated and will change in a future version. Call result.infer\_objects(copy=False) instead. To opt-in to the future behavior, set  
`pd.set\_option('future.no\_silent\_downcasting', True)`  
df[col]=df[col].fillna(df[col].mode()[0])

```
[9]: df.isna().sum()
```

```
[9]: id          0
age           0
sex           0
dataset       0
cp            0
trestbps      0
chol          0
fbs           0
restecg       0
thalch        0
exang         0
oldpeak       0
slope         0
ca            0
thal          0
num           0
dtype: int64
```

```
[10]: df1=df[['age','cp','chol','thalch']]
df2=df[['exang','slope','num']]
```

```
[11]: merged_df=pd.concat([df1,df2],axis=1)
merged_df
df.head()
```

```
[11]:
```

	id	age	sex	dataset	cp	trestbps	chol	fbs	\
0	1	63	Male	Cleveland	typical angina	145.0	233.0	True	
1	2	67	Male	Cleveland	asymptomatic	160.0	286.0	False	
2	3	67	Male	Cleveland	asymptomatic	120.0	229.0	False	
3	4	37	Male	Cleveland	non-anginal	130.0	250.0	False	
4	5	41	Female	Cleveland	atypical angina	130.0	204.0	False	

		restecg	thalch	exang	oldpeak	slope	ca	\
0	lv hypertrophy	150.0	False	2.3	downsloping	0.0		
1	lv hypertrophy	108.0	True	1.5	flat	3.0		

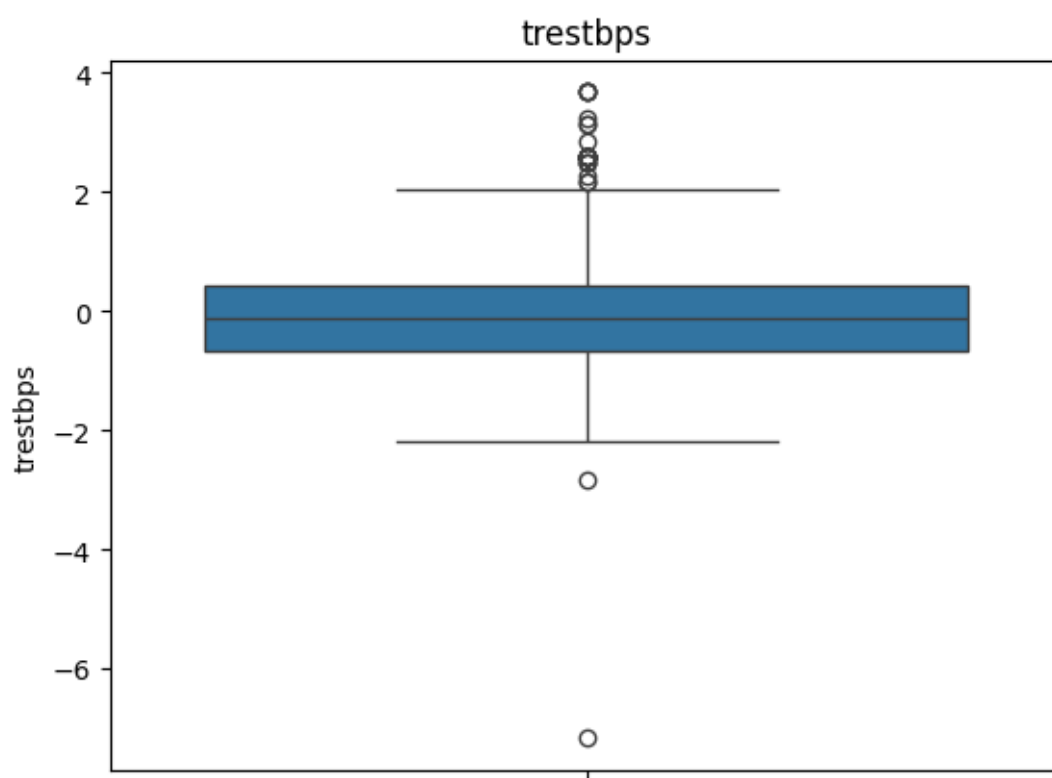

2	lv hypertrophy	129.0	True	2.6	flat	2.0
3	normal	187.0	False	3.5	downsloping	0.0
4	lv hypertrophy	172.0	False	1.4	upsloping	0.0


	thal	num
0	fixed defect	0
1	normal	2
2	reversable defect	1
3	normal	0
4	normal	0



```
[12]: from sklearn.preprocessing import StandardScaler, LabelEncoder
label_encoder=LabelEncoder()
scaler=StandardScaler()
cat_col=['sex', 'dataset', 'cp', 'restecg', 'slope', 'thal']
num_cols=['age', 'trestbps', 'chol', 'thalch', 'oldpeak', 'ca']
```


```
[13]: for col in cat_col:
        df[col]=label_encoder.fit_transform(df[col].astype(str))
df[num_cols]=scaler.fit_transform(df[num_cols])

for col in num_cols:
    sns.boxplot(df[col])
    plt.title(col)
    plt.show()
```







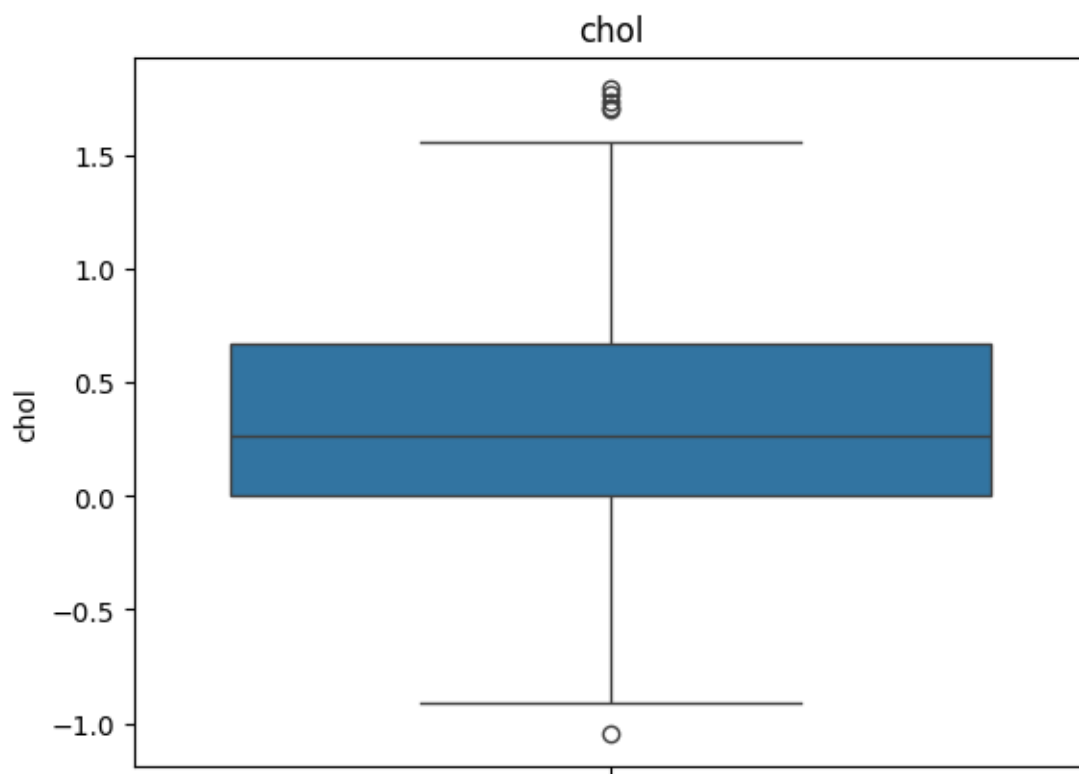







```
[14]: def remove_outliers(df, col):  
    Q1=df[col].quantile(0.25)  
    Q3=df[col].quantile(0.75)  
    IQR=Q3-Q1  
    lower= Q1 - 1.5 * IQR  
    upper= Q3 + 1.5 * IQR  
    return df[(df[col]>=lower) & (df[col]<=upper) ]  
  
for col in num_cols:  
    df=remove_outliers(df, col)  
  
for col in num_cols:  
    sns.boxplot(df[col])  
    plt.title(col)  
    plt.show()
```











```
[15]: df.head()
```

```
[15]:
```

	id	age	sex	dataset	cp	trestbps	chol	fbs	restecg	\
166	167	-0.160397	1	0	2	0.318305	0.219192	False	1	
192	193	-1.115855	1	0	0	-0.007183	0.439581	True	0	
287	288	0.476575	1	0	1	-0.386918	0.191643	False	1	
302	303	-1.646666	1	0	2	0.318305	-0.221586	False	1	
303	304	-2.708286	1	0	1	-0.115679	-0.616449	False	0	

	thalch	exang	oldpeak	slope	ca	thal	num
166	1.251922	False	-0.834397	2	-2.050756e-16	1	0
192	0.217089	True	-0.739448	1	-2.050756e-16	2	1
287	0.256891	False	-0.454603	1	-2.050756e-16	2	0
302	1.411127	False	-0.834397	2	-2.050756e-16	1	0
303	1.888743	False	-0.834397	1	-2.050756e-16	1	0

```
[16]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split

y=df['num']
```

```
X=df.drop(['id','num'],axis=1)

X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.
↪2,random_state=42)

from sklearn.metrics import accuracy_score,classification_report

model=LogisticRegression()
model.fit(X_train,y_train)
y_pred=model.predict(X_test)
```

C:\Users\AMOL\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\linear\_model\\_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
n\_iter\_i = \_check\_optimize\_result(

```
[17]: print(f'Accuracy:',accuracy_score(y_test,y_pred))
      print(f'Classification Report:',classification_report(y_test,y_pred))
```

Accuracy: 0.75

Classification Report:			precision	recall	f1-score	support
0	0.87	0.92	0.90	51		
1	0.60	0.60	0.60	20		
2	1.00	0.50	0.67	6		
3	0.14	0.20	0.17	5		
4	0.00	0.00	0.00	2		
accuracy			0.75	84		
macro avg			0.52	0.44	0.47	84
weighted avg			0.75	0.75	0.74	84

C:\Users\AMOL\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\metrics\\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
C:\Users\AMOL\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\metrics\\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted

```
samples. Use `zero_division` parameter to control this behavior.  
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
C:\Users\AMOL\AppData\Local\Programs\Python\Python313\Lib\site-  
packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning:  
Precision is ill-defined and being set to 0.0 in labels with no predicted  
samples. Use `zero_division` parameter to control this behavior.  
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
[ ]: new_data = [[-0.2, 1, 0, 2, 0.25, 0.3, 1, 1, 1.2, 0, 0.0, 2, 1, 2]]
```